

Predicting house prices: a regression example

Anton Antonov
MathematicaVsR at GitHub
May 2018

Introduction

This notebook is part of the MathematicaVsR at GitHub project “DeepLearningExamples”.

This notebook has code that corresponds to code in the book “Deep learning with R” by F. Chollet and J. J. Allaire. See the GitHub repository: <https://github.com/jjallaire/deep-learning-with-r-notebooks> ; specifically the notebook “Predicting house prices: a regression example”.

Get code

Here we load a package with utilities:

```
In[1]:= Import["https://raw.githubusercontent.com/antononcube/MathematicaForPrediction/master/MathematicaForPredictionUtilities.m"]  
» Importing from GitHub: MosaicPlot.m  
» Importing from GitHub: CrossTabulate.m
```

Get data

We can get the data from Wolfram’s data repository:

```
In[2]:= (*ResourceObject["Sample Data: Boston Homes"]*)
```

or we can directly use ExampleData:

```
In[3]:= trainingData = ExampleData[{"MachineLearning", "BostonHomes"}, "TrainingData"];  
testData = ExampleData[{"MachineLearning", "BostonHomes"}, "TestData"];
```

Here is a description of the variable names:

```
In[5]:= colNames = Most@Flatten[List@@ExampleData[{"MachineLearning", "BostonHomes"}, "VariableDescriptions"]];
GridTableForm[List /@ colNames]
```

Out[6]=

#	l
1	Per capita crime rate by town
2	Proportion of residential land zoned for lots over 25000 square feet
3	Proportion of non-retail business acres per town
4	Charles River dummy variable (1 if tract bounds river, 0 otherwise)
5	Nitrogen oxide concentration (parts per 10 million)
6	Average number of rooms per dwelling
7	Proportion of owner-occupied units built prior to 1940
8	Weighted mean of distances to five Boston employment centers
9	Index of accessibility to radial highways
10	Full-value property-tax rater per \$10000
11	Pupil-teacher ratio by town
12	1000(Bk-0.63)^2 where Bk is the proportion of blacks by town
13	Lower status of the population (percent)
14	Median value of owner-occupied homes in \$1000s

Data summaries

```
In[7]:= Dimensions[Flatten@*List@@@trainingData]
```

```
Out[7]= {338, 14}
```

```
In[8]:= Dimensions[Flatten@*List@@@testData]
```

```
Out[8]= {168, 14}
```

```
In[9]:= Magnify[RecordsSummary[trainingData, Thread → True], 0.6]
```

Out[9]=

1 column 1	2 column 2	3 column 3	4 column 4	5 column 5	6 column 6	7 column 7	8 column 8	9 column 9	10 column 10	11 column 11	12 column 12	13 column 13	1 column 1
Min 0.00906	1st Qu 0	Min 0.46	1st Qu 0	Min 0.385	Min 3.561	Min 2.9	Min 1.137	Min 1	Min 187	Min 12.6	Min 2.52	Min 1.73	Min 5
1st Qu 0.08873	Median 0	1st Qu 5.19	3rd Qu 0	1st Qu 0.453	1st Qu 5.888	1st Qu 45	1st Qu 2.1099	1st Qu 4	1st Qu 277	1st Qu 17.4	Mean 358.954	1st Qu 6.93	1st Qu 17.2
Median 0.274475	Min 0	Median 8.56	Median 0	Median 0.538	Median 6.209	Mean 69.1896	Median 3.2157	Median 5	Median 315	Mean 18.4808	1st Qu 376.14	Median 11.645	Median 21.3
3rd Qu 2.37857	Mean 10.9734	Mean 10.8893	Min 0	Mean 0.551747	Mean 6.31367	Median 79.05	Mean 3.786	3rd Qu 8	Mean 400.757	Median 19.1	Median 391.52	Mean 12.5943	Mean 22.7145
Mean 3.50511	3rd Qu 12.5	3rd Qu 18.1	Mean 0.0680473	3rd Qu 0.614	3rd Qu 6.683	3rd Qu 94	3rd Qu 4.9671	Mean 9.24556	3rd Qu 666	3rd Qu 20.2	3rd Qu 396.33	3rd Qu 16.96	3rd Qu 26.4
Max 88.9762	Max 100	Max 27.74	Max 1	Max 0.871	Max 8.725	Max 100	Max 12.1265	Max 24	Max 711	Max 22	Max 396.9	Max 37.97	Max 50

```
In[10]:= Magnify[RecordsSummary[testData, Thread → True], 0.6]
```

Out[10]=

1 column 1	2 column 2	3 column 3	4 column 4	5 column 5	6 column 6	7 column 7	8 column 8	9 column 9	10 column 10	11 column 11	12 column 12	13 column 13	1 column 1
Min 0.00632	1st Qu 0	Min 1.21	1st Qu 0	Min 0.394	Min 3.863	Min 6.2	Min 1.1296	Min 1	Min 188	Min 12.6	Min 0.32	Min 2.88	Min 8.4
1st Qu 0.068935	Median 0	1st Qu 4.93	3rd Qu 0	1st Qu 0.446	1st Qu 5.8765	1st Qu 45.35	1st Qu 1.96175	1st Qu 4	1st Qu 285.5	1st Qu 16.6	Mean 352.088	1st Qu 7.06	1st Qu 16.1
Median 0.21693	Min 0	Median 10.59	Median 0	Median 0.538	Median 6.2005	Mean 67.3381	Median 3.1423	Median 5	Median 377	Mean 18.4048	1st Qu 370.5	Median 11.25	Median 20.95
Mean 3.83163	Mean 12.1488	Mean 11.6348	Min 0	Mean 0.560626	Mean 6.22622	Median 72.95	Mean 3.81324	Mean 10.1607	Mean 423.286	Median 18.95	Median 391.44	Mean 12.7714	Mean 22.1673
3rd Qu 4.75225	3rd Qu 12.5	3rd Qu 18.1	Mean 0.0714286	3rd Qu 0.668	3rd Qu 6.543	3rd Qu 94.25	3rd Qu 5.4007	3rd Qu 24	3rd Qu 666	3rd Qu 20.2	3rd Qu 395.69	3rd Qu 16.92	3rd Qu 24.75
Max 73.5341	Max 95	Max 27.74	Max 1	Max 0.871	Max 8.78	Max 100	Max 10.7103	Max 24	Max 711	Max 21.2	Max 396.9	Max 34.77	Max 50

Standardize data

First we turn the data sets into arrays:

```
In[12]:= trainingData = Flatten@*List@@@trainingData;
```

```
testData = Flatten@*List@@@testData;
```

Note that the standardisation of the data is done by using means and standard deviations derived from the training data only.

```
In[14]:= means = Mean[trainingData];
```

```
sds = N[StandardDeviation /@ Transpose[trainingData]];
```


The neural network

Build the network:

```
In[21]:= model =  
  NetChain[{  
    LinearLayer[64, "Input" → Dimensions[trainingData[[All, 1]] [[2]]],  
    ElementwiseLayer[Ramp],  
    LinearLayer[64],  
    ElementwiseLayer[Ramp],  
    LinearLayer[1]  
  }, "Output" → "Scalar"]
```

Out[21]= NetChain[

uninitialized

Input

vector (size: 13)

1 LinearLayer

vector (size: 64)

2 Ramp

vector (size: 64)

3 LinearLayer

vector (size: 64)

4 Ramp

vector (size: 64)

5 LinearLayer

vector (size: 1)

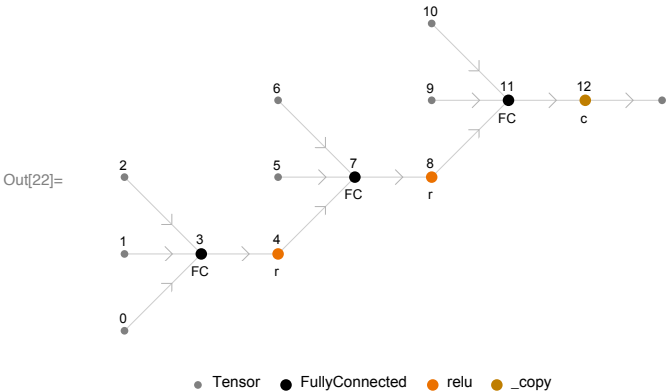
Output

scalar

]

Visualize the network:

```
In[22]:= NetInformation[model, "MXNetNodeGraphPlot"]
```



Train the network:

```
In[23]:= tNet = NetTrain[model, trainingData,  
  ValidationSet → Scaled[.05], "MaxTrainingRounds" → 80]
```

Out[23]= NetChain[

Input

vector (size: 13)

1 LinearLayer

vector (size: 64)

2 Ramp

vector (size: 64)

3 LinearLayer

vector (size: 64)

4 Ramp

vector (size: 64)

5 LinearLayer

vector (size: 1)

Output

scalar

]

Testing

Predict results for the test data:

```
In[32]:= pRes = tNet /@ testData[[All, 1]];
Short[pRes]
```

```
Out[33]/Short= {0.588873, 0.909774, <<164>>, -0.952313, -0.37624}
```

Mean absolute error:

```
In[34]:= Mean[Abs[pRes - (Last /@ testData)]]
```

```
Out[34]= 0.26433
```

Root mean square error:

```
In[35]:= RootMeanSquare[pRes - (Last /@ testData)]
```

```
Out[35]= 0.451546
```

Compare the predicted results with the actual values using a scatter plot:

```
In[36]:= Show[
  ListPlot[Transpose[{pRes, Last /@ testData}], FrameLabel -> {"Predicted", "Actual"},
  PlotTheme -> "Detailed", PlotRange -> All, ImageSize -> Large],
  Plot[{x, x}, {x, Min[pRes], Max[pRes]}]]
```

